

COMPOSE YOURSELF: AVERAGE-VELOCITY FLOW MATCHING FOR ONE-STEP SPEECH ENHANCEMENT

Gang Yang[†] Yue Lei[†] Wenxin Tai[†] Jin Wu[†] Jia Chen[†] Ting Zhong[†] Fan Zhou^{†*}

[†] University of Electronic Science and Technology of China, Chengdu, Sichuan, China

ABSTRACT

Diffusion and flow matching (FM) models have achieved remarkable progress in speech enhancement (SE), yet their dependence on multi-step generation is computationally expensive and vulnerable to discretization errors. Recent advances in one-step generative modeling, particularly MeanFlow, provide a promising alternative by reformulating dynamics through average velocity fields. In this work, we present COSE, a one-step FM framework tailored for SE. To address the high training overhead of Jacobian-vector product (JVP) computations in MeanFlow, we introduce a velocity composition identity to compute average velocity efficiently, eliminating expensive computation while preserving theoretical consistency and achieving competitive enhancement quality. Extensive experiments on standard benchmarks show that COSE delivers up to 5x faster sampling and reduces training cost by 40%, all without compromising speech quality. Code is available at <https://github.com/ICDM-UESTC/COSE>.

Index Terms— generative model, speech enhancement, flow matching, average velocity

1. INTRODUCTION

Speech enhancement (SE) aims to restore clean speech signals from recordings corrupted by noise, reverberation, and encoding artifacts [1]. It benefits both human perception and downstream tasks such as automatic speech recognition (ASR) [2]. Traditional methods based on statistical signal processing have been widely studied, but they often struggle to handle non-stationary or complex noise. The advent of deep learning has since driven significant progress by learning to directly map noisy speech to its clean counterpart [3]. Recently, generative models [4–6] have gained prominence, excelling at capturing complex speech distributions and offering a powerful approach that preserves both speech quality and intelligibility.

Among generative models, diffusion and flow matching (FM) approaches have emerged as promising solutions for speech enhancement, demonstrating strong robustness to unseen noise [6] and are considered promising solutions for speech enhancement [7–11]. Diffusion models use a forward stochastic differential equation (SDE) to transform clean speech, FM models directly learn a deterministic velocity field defined by an ordinary differential equation (ODE). Despite their impressive performance, these models still face challenges. Errors from large-step discretization degrade few-step performance, and both approaches remain computationally costly, often requiring five or more function evaluations (NFEs) during sampling [12–14].

Currently, a growing body of work has attempted to bridge this efficiency gap through one-step generation techniques in FM [15, 16]. These methods avoid the error accumulation inherent in multi-step

generation [17] and are considered promising for improving computational efficiency. Among these, the recently proposed MeanFlow framework [16] is particularly notable. It provides an elegant mathematical formulation of generative dynamics by introducing the concept of average velocity over a time interval. This approach enables direct one-step generation via learned velocity fields, which holds great promise for multi-step generative speech enhancement by significantly improving the efficiency of generation.

In this work, we propose COSE (*Compose velOcity in Speech Enhancement*), a framework that integrates Meanflow [16] into SE, enabling efficient, one-step generation. Further, by leveraging the properties of ODEs, we compute the average velocity via the velocity composition identity, which avoids Jacobian–vector product (JVP) calculations and substantially reduces training overhead. Specifically, we incorporate Meanflow into SE by modeling the average velocity between two points on a curved trajectory, enabling the model to generate clean speech in one step. Building on this, we exploit the uniqueness property of ODEs to decompose displacements into compositions of two segment velocities, thereby eliminating the extra cost of direct JVP computations. Moreover, we show that velocity composition yields a solution equivalent to other one-step generation methods, consistent with the self-consistency principle underlying the trajectory properties of flow matching.

Our main contributions are: (1) We introduce COSE, a novel framework that integrates Meanflow into speech enhancement to enable efficient one-step generation. (2) We analyze that JVP computation incurs significant overhead. Instead, we leveraged ODE properties to calculate average velocity, effectively avoiding JVP and its associated computational cost. (3) Our extensive experiments on standard benchmarks demonstrate that COSE achieves at least 5x faster sampling and reduces GPU memory and training time by 40% compared to MeanFlow, while maintaining equivalent performance.

2. PRELIMINARY

2.1. Speech Enhancement

Speech enhancement (SE) improves intelligibility and quality by suppressing noise while preserving natural speech characteristics. In the single-channel setting, the noisy signal can be expressed as:

$$\mathbf{y} = \mathbf{x} + \mathbf{n}, \quad (1)$$

where \mathbf{x} is the clean speech and \mathbf{n} is additive noise. The objective is to estimate \mathbf{x} from \mathbf{y} .

2.2. Flow Matching for Speech Enhancement

Flow Matching (FM) models generative modeling as the learning of continuous-time velocity fields. It aims to learn a continuous-time trajectory $\mathbf{x}_t \in \mathbb{R}^d$, where $t \in [0, 1]$ from a simple distribution to a

*Corresponding author.

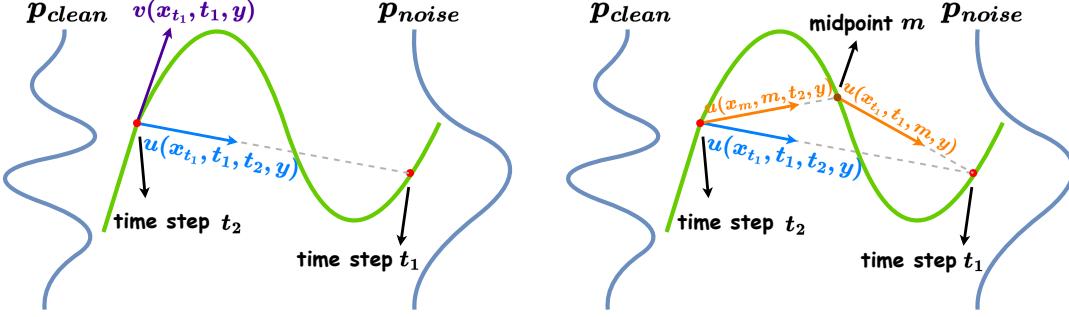


Fig. 1: Overview of velocity modeling. **Left:** Comparison between instantaneous and average velocity. The green line represents a generative trajectory between clean speech distribution and Gaussian noise distribution. **Right:** An example of velocity composition identity. For a midpoint m between time steps t_1 and t_2 , the two sub-average velocities can synthesize the average velocity of the large step.

complex clean speech distribution. A linear interpolation trajectory is commonly employed to construct intermediate states between a clean speech \mathbf{x}_0 and a sample from the prior $\mathbf{x}_1 \sim \mathcal{N}(0, \mathbf{I})$ [18]:

$$\mathbf{x}_t = (1 - t)\mathbf{x}_0 + t\mathbf{x}_1, \quad \mathbf{x}_1 \sim \mathcal{N}(0, \mathbf{I}), \quad (2)$$

which defines a conditional distribution $p_t(\mathbf{x}_t \mid \mathbf{x}_0)$ from which intermediate states can be sampled explicitly. The evolution of \mathbf{x}_t is governed by a time-dependent instantaneous velocity field $v_t(\mathbf{x}_t)$, described by the following ordinary differential equation (ODE):

$$\frac{d\mathbf{x}_t}{dt} = v_t(\mathbf{x}_t), \quad t \in [0, 1]. \quad (3)$$

Then a neural network $v_\theta(\mathbf{x}_t, t, \mathbf{y})$ is trained to approximate the velocity field with the Conditional Flow Matching (CFM) objective [19]:

$$\mathcal{L}_{\text{CFM}}(\theta) = \mathbb{E}_{t \sim \mathcal{U}(0, 1), \mathbf{x}_0, \mathbf{x}_1 \sim \mathcal{N}(0, \mathbf{I})} \|v_\theta(\mathbf{x}_t, t, \mathbf{y}) - (\mathbf{x}_1 - \mathbf{x}_0)\|^2, \quad (4)$$

where \mathbf{x}_t is sampled from eq. (2) and \mathbf{y} represents the noisy speech.

During inference, eq. (3) is solved numerically using the Euler method with trained $v_\theta(\mathbf{x}_t, t)$ to generate clean speech:

$$\mathbf{x}_{t-\Delta t} = \mathbf{x}_t - \Delta t v_\theta(\mathbf{x}_t, t, \mathbf{y}), \quad \forall t \in [0, 1], \quad (5)$$

where $\Delta t = 1/N$, and N represents the number of function evaluations (NFEs). By iteratively solving eq. (5), the speech can be refined from the noisy version to the clean version.

3. METHODOLOGY

In this section, we extend MeanFlow to speech enhancement and introduce a velocity composition identity to avoid computationally expensive operations. A overview of the velocity modeling approach is illustrated in fig. 1.

3.1. MeanFlow for Speech Enhancement

Some recent work [15, 16] has observed that although the interpolated trajectories are designed to be straight during training, FM models can only predict the marginal velocity based on the current state during inference, which will naturally cause the trajectory to bend and deviate from the original straight line, leading one-step generation to fail completely.

To address this challenge, MeanFlow builds an explicit model of the average velocity between two points on curved trajectory [16], as

illustrated in fig. 1 left. Specifically, we first define the displacement $k(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$ between two time steps t_1 and t_2 as the integral of the instantaneous velocity field over that time interval as:

$$k(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) = \int_{t_2}^{t_1} v(\mathbf{x}_\tau, \tau, \mathbf{y}) d\tau, \quad (6)$$

where \mathbf{x}_{t_1} represents intermediate states at time step t , \mathbf{y} represents noisy speech, and $v(\mathbf{x}_\tau, \tau, \mathbf{y})$ represents instantaneous velocity. The average velocity $u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$ is then defined as this displacement divided by the time interval:

$$u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) \triangleq \frac{k(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})}{t_1 - t_2} = \frac{1}{t_1 - t_2} \int_{t_2}^{t_1} v(\mathbf{x}_\tau, \tau, \mathbf{y}) d\tau. \quad (7)$$

By differentiating both sides of the equation with respect to t_1 simultaneously, this identity can be transformed into:

$$u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) = v(\mathbf{x}_t, t_1, \mathbf{y}) - (t_1 - t_2) \frac{d}{dt_1} u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}), \quad (8)$$

where $\frac{d}{dt_1} u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$ denotes the total derivative of the average velocity function u with respect to time t_1 .

According to eq. (8), we can train a neural network $u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$ to directly model this average velocity field. The training objective encourages the network to satisfy the MeanFlow Identity:

$$\mathcal{L}(\theta) = \mathbb{E}_{\mathbf{x}_0, \mathbf{y}, t_1, t_2} \|u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) - \text{sg}(u_{\text{tgt}})\|_2^2, \quad (9)$$

where $u_{\text{tgt}} = v(\mathbf{x}_{t_1}, t_1, \mathbf{y}) - (t_1 - t_2) \frac{d}{dt_1} u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$. Following [16], we adopt the stop-gradient operation, denoted as $\text{sg}(\cdot)$, and replace $v(\mathbf{x}_{t_1}, t_1, \mathbf{y})$ with $\mathbf{x}_1 - \mathbf{x}_0$.

In eq. (9), $\frac{d}{dt_1} u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$ is computed as the Jacobian–vector product (JVP) between $[\partial_{\mathbf{x}_t} u, \partial_{t_1} u, \partial_{t_2} u, \partial_y u]$ and the tangent vector $[v, 1, 0, 0]$, which can be implemented via PyTorch’s JVP function. However, MeanFlow’s reliance on JVP incurs significant overhead, limiting training efficiency.

3.2. Revisiting ODE Properties for Efficient Training

Computational Overhead of JVP. The JVP computes the directional derivative of a function along a given tangent vector. Modern automatic differentiation frameworks, such as PyTorch [20], use dual numbers to efficiently implement it as:

$$u(\mathbf{x}_{t_1} + \mathbf{v}\delta, t_1 + \delta, t_2, \mathbf{y}) = u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) + \left(\frac{\partial u}{\partial \mathbf{x}_{t_1}} \cdot \mathbf{v} + \frac{\partial u}{\partial t_1} \right) \delta,$$

where δ is an infinitesimal for differentiation, and $u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$ corresponds to the forward output, and the second term corresponds to the *JVP* result. This approach requires only a single forward pass and theoretically introduces a slight computational cost. However, in practice, maintaining dual-numbered computations incurs memory allocations, additional arithmetic operations, and expanded computation graphs, leading to noticeable computational and memory costs [21]. Furthermore, *JVP* implementations differ across frameworks (e.g., PyTorch vs. JAX), adding development complexity and reducing portability. To address these limitations, we revisit the Mean-Flow training method and propose a general solution that entirely circumvents *JVP*, achieving efficient and memory-friendly training.

Velocity Composition of ODE. Inspired by previous work [15, 17], we derive a more general training framework for one-step generative models from the fundamental perspective of the uniqueness of ODE solutions. This framework allows us to bypass costly *JVP* computations. Specifically, the ODE describes the deterministic evolution of a system. In line with previous studies [19, 22], we assume that the instantaneous velocity field is locally Lipschitz continuous in the region traversed by the generated trajectories, denoted as \mathcal{X}_t :

$$\forall \mathbf{x}_1, \mathbf{x}_2 \in \mathcal{X}_t, \quad \|v(\mathbf{x}_1, \tau, \mathbf{y}) - v(\mathbf{x}_2, \tau, \mathbf{y})\| \leq L \|\mathbf{x}_1 - \mathbf{x}_2\|, \quad (10)$$

where \mathcal{X}_t represents the set of states visited by the ODE solution during generation. Under this Lipschitz condition, the ODE satisfies the **semigroup property** [23]: its evolution from time t_2 to t_1 can be decomposed into a two-step process, first from t_2 to an intermediate time m , and then from m to t_1 . This can be formalized as:

$$\Phi_{t_2 \rightarrow t_1} = \Phi_{m \rightarrow t_1} \circ \Phi_{t_2 \rightarrow m}, \quad s = t_2 + \alpha(t_1 - t_2), \quad (11)$$

where $\Phi_{m \rightarrow t_1}$ is the evolution operator, namely $\Phi_{t_2 \rightarrow t_1}(\mathbf{x}_{t_2}) = \mathbf{x}_{t_1}$ and $\alpha \in [0, 1]$, which maps the state at time t_2 to its corresponding state at time t_1 along the ODE trajectory. Consequently, the total displacement can be expressed as the sum of two sub-displacements:

$$k(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) = k(\mathbf{x}_{t_1}, t_1, m, \mathbf{y}) + k(\mathbf{x}_m, m, t_2, \mathbf{y}), \quad (12)$$

According to the definition in eq. (7) and eq. (12), we can further transform them into the following interpolation identity:

$$u(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) = u(\mathbf{x}_m, m, t_2, \mathbf{y}) + \alpha(u(\mathbf{x}_{t_1}, t_1, m, \mathbf{y}) - u(\mathbf{x}_m, m, t_2, \mathbf{y})), \quad (13)$$

which can be interpreted as the velocity composition identity, combining two velocities into one, as shown in fig. 1 right. Here, the parameter α is randomly sampled from $[0, 1]$, encouraging the model to learn consistency across different temporal granularities during training. Based on eq. (13), we train the loss function of COSE without requiring expensive *JVP* operations as:

$$\mathcal{L}_{\text{COSE}}(\theta) = \mathbb{E}_{t, \mathbf{x}_t, \mathbf{y}} \|u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) - \text{sg}(u_{\text{tgt}})\|_2^2, \quad (14)$$

where $\text{sg}(\cdot)$ denotes the stop-gradient operator following [16] and u_{tgt} defined as:

$$u_{\text{tgt}} = u_\theta(\mathbf{x}_m, m, t_2, \mathbf{y}) + \alpha(u_\theta(\mathbf{x}_{t_1}, t_1, m, \mathbf{y}) - u_\theta(\mathbf{x}_m, m, t_2, \mathbf{y})).$$

The overall training procedure is summarized in algorithm 1.

Beyond this, we establish a connection between velocity composition identity and self-consistency loss functions, ensuring the learning of a self-consistent and unidirectional average velocity field. Specifically, we divide the time interval $[t_2, t_1]$ into two steps based on intermediate point m and let $d_1 = t_1 - m$ and $d_2 = m - t_2$

Algorithm 1 COSE Training Algorithm

Require: Neural network u_θ , a batch of clean data x and noisy data \mathbf{y} , optimizer.

- 1: Sample time points t_2, t_1 such that $0 \leq t_2 \leq t_1 \leq 1$.
- 2: Sample $\alpha \sim \mathcal{U}(0, 1)$, set $m = t_1 + \alpha(t_2 - t_1)$.
- 3: Sample prior $\epsilon \sim \mathcal{N}(0, I)$.
- 4: Construct flow path point at time t : $\mathbf{x}_{t_1} = (1 - t_1)x + t_1\epsilon$.
- 5: $u_2 = u_\theta(\mathbf{x}_{t_1}, t_1, m, \mathbf{y})$, $\mathbf{x}_m = \mathbf{x}_{t_1} - (t_1 - m)u_2$.
- 6: $u_1 = u_\theta(\mathbf{x}_m, m, t_2, \mathbf{y})$, $\mathbf{x}_{t_2} = \mathbf{x}_m - (m - t_2)u_1$.
- 7: $u_{t_1 t_2} = u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y})$, $u_{\text{tgt}} = u_1 + \alpha(u_2 - u_1)$.
- 8: Update θ via $\nabla_\theta \mathcal{L}_{\text{COSE}}$, $\mathcal{L}_{\text{COSE}} = \|u_{t_1 t_2} - \text{sg}(u_{\text{tgt}})\|_2^2$

represent the time intervals between t_1 , m , and t_2 . Based on this, eq. (13) can be transformed into:

$$u_\theta(\mathbf{x}_{t_1}, t_1, t_2, \mathbf{y}) = \alpha \cdot u(\mathbf{x}_{t_1}, t_1, t_1 - d_1, \mathbf{y}) + (1 - \alpha) \cdot u(\mathbf{x}_{t_1 - d_1}, t_1 - d_1 - d_2, \mathbf{y}), \quad (15)$$

where $\mathbf{x}_{t_1 - d_1} = \mathbf{x}_{t_1} - u_\theta(\mathbf{x}_{t_1}, t_1, d_1, \mathbf{y})$. We observe that eq. (15) is algebraically equivalent to the self-consistency identity in [15, eq. (4)], which enables propagating generation across step scales (multi→few→one) and joint end-to-end training [17]. COSE builds on this principle to provide a flexible and general framework that accommodates arbitrary time intervals and diverse evolution patterns.

4. EXPERIMENTS

4.1. Experimental Settings

Datasets. Our method is evaluated on the VoiceBank-DEMAND (VBD) [25] and CHiME-4 [26] datasets. The VBD training set contains recordings from 28 speakers (14 female, 14 male) with 8 noise types. For testing, we use the VBD test set with two unseen speakers and two new noise types. For a robust cross-dataset evaluation, we follow [6, 11] and use the fifth-microphone recordings from CHiME-4, which feature real-world noise from four environments: street, walking area, cafeteria, and bus.

Implementation Details. Following previous work [16], we extend the NCSN++ architecture with a Fourier-embedded channel and concatenated (t_1, t_2) steps to model average velocity, with negligible parameter overhead. We also adopt the same training setup: sampling timestep from a lognorm($-0.4, 1.0$), setting $t_2 = t_1$ with 50% probability, and using the adaptive L2 loss $\mathcal{L} = w|\Delta|_2^{2\gamma}$, where $w = 1/(|\Delta|_2^2 + c)^p$ and $c = 10^{-3}$.

Metrics. We quantitatively evaluate performance using several standard metrics: Perceptual Evaluation of Speech Quality (PESQ) [27] for perceived quality, Extended Short-Time Objective Intelligibility (ESTOI) [28] for intelligibility, and the scale-invariant measures SI-SDR, SI-SIR, and SI-SAR [29].

Configurations. We follow [8] in using the same backbone and audio preprocessing, applying a 512-sample STFT with a 128-sample hop to obtain 256 frequency bins. Models are trained for 200 epochs on a single NVIDIA 4090 GPU with a learning rate of 5×10^{-5} .

Baselines. We compare our proposed methodology with three baseline models: SGMSE+ [8], StoRM [24], VPIDM [9], FlowSE [13, 30], LARF [14]. For FlowSE, we adopt the method from two concurrent conference works [13, 30] and re-implement it following the diffusion schedule and architecture in [13, 19], as we did not reproduce the reported results in our setting using the available implementation. Following prior work [9], we also incorporate the NCSN++ [8] architecture as a discriminative model in our comparisons.

Table 1: Speech enhancement results on VoiceBank-DEMAND and ChiME-4 datasets. All experimental results are presented as mean \pm standard deviation. The best results are highlighted in bold.

Method	NFE	VoiceBank- DEMAND					CHiME-4				
		PESQ	ESTOI	SI-SDR	SI-SIR	SI-SAR	PESQ	ESTOI	SI-SDR	SI-SIR	SI-SAR
Mixture	-	1.97 \pm 0.75	0.79 \pm 0.15	8.4 \pm 5.6	8.5 \pm 5.6	47.5 \pm 10.4	1.27 \pm 0.16	0.68 \pm 0.08	7.5 \pm 2.1	7.5 \pm 2.1	46.7 \pm 10.3
NCSN++ [8]	1	2.87 \pm 0.74	0.87 \pm 0.10	19.1 \pm 3.5	31.5 \pm 7.2	20.0 \pm 3.5	1.27 \pm 0.16	0.68 \pm 0.08	7.4 \pm 2.1	7.4 \pm 2.1	42.4 \pm 5.4
SGMSE+ [8]	15	2.80 \pm 0.58	0.86 \pm 0.10	17.2 \pm 3.6	26.9 \pm 5.2	17.9 \pm 3.5	1.82 \pm 0.29	0.82 \pm 0.07	13.5 \pm 2.3	25.2 \pm 3.4	13.9 \pm 2.2
SGMSE+ [8]	5	1.15 \pm 0.08	0.63 \pm 0.10	7.9 \pm 1.8	18.6 \pm 4.1	8.4 \pm 1.6	1.09 \pm 0.04	0.58 \pm 0.09	5.3 \pm 1.6	15.2 \pm 2.1	5.8 \pm 1.6
SGMSE+ [8]	1	1.06 \pm 0.10	0.02 \pm 0.02	-25.3 \pm 1.7	23.3 \pm 9.9	-25.3 \pm 1.7	1.08 \pm 0.14	0.02 \pm 0.02	-26.6 \pm 1.2	22.8 \pm 8.9	-26.6 \pm 1.2
StoRM [24]	15	2.77 \pm 0.57	0.87 \pm 0.10	18.5 \pm 3.3	30.9 \pm 6.7	19.1 \pm 3.3	1.82 \pm 0.30	0.84 \pm 0.06	14.5 \pm 2.2	26.0 \pm 3.8	14.8 \pm 2.2
StoRM [24]	5	1.25 \pm 0.09	0.70 \pm 0.10	11.5 \pm 1.4	29.8 \pm 6.1	11.6 \pm 1.4	1.19 \pm 0.08	0.72 \pm 0.08	9.9 \pm 1.6	24.5 \pm 3.3	10.0 \pm 1.6
StoRM [24]	1	1.04 \pm 0.02	0.10 \pm 0.03	-16.9 \pm 1.5	22.2 \pm 7.0	-16.9 \pm 1.5	1.05 \pm 0.12	0.10 \pm 0.03	-17.7 \pm 1.1	18.3 \pm 3.4	-17.7 \pm 1.1
VPIDM [9]	15	2.92 \pm 0.61	0.87 \pm 0.10	18.8 \pm 2.5	28.4 \pm 5.3	19.4 \pm 1.9	1.67 \pm 0.20	0.80 \pm 0.07	14.2 \pm 1.9	23.2 \pm 2.9	13.5 \pm 2.9
VPIDM [9]	5	1.85 \pm 0.12	0.83 \pm 0.10	14.6 \pm 2.1	17.7 \pm 3.9	14.9 \pm 1.6	1.03 \pm 0.01	0.33 \pm 0.07	-6.7 \pm 1.5	12.2 \pm 1.8	-6.6 \pm 1.5
VPIDM [9]	1	1.04 \pm 0.01	0.29 \pm 0.06	-7.2 \pm 1.5	14.3 \pm 3.7	-7.2 \pm 1.5	1.04 \pm 0.01	0.29 \pm 0.06	-7.2 \pm 1.5	14.3 \pm 3.7	-7.2 \pm 1.5
FlowSE [19]	5	2.96 \pm 0.73	0.87 \pm 0.10	18.8 \pm 3.5	31.7 \pm 7.0	19.8 \pm 3.5	1.73 \pm 0.31	0.84 \pm 0.06	14.2 \pm 2.3	24.3 \pm 3.5	14.7 \pm 2.3
FlowSE [19]	1	1.37 \pm 0.15	0.74 \pm 0.10	12.1 \pm 2.0	22.2 \pm 4.6	12.7 \pm 1.7	1.06 \pm 0.03	0.53 \pm 0.09	2.0 \pm 1.5	16.1 \pm 2.3	2.1 \pm 1.5
LARF [14]	5	2.98 \pm 0.75	0.87 \pm 0.10	18.8 \pm 3.4	26.6 \pm 7.4	20.2 \pm 3.5	1.67 \pm 0.27	0.84 \pm 0.07	14.8 \pm 2.4	25.7 \pm 3.8	15.2 \pm 2.3
LARF [14]	1	2.97 \pm 0.70	0.87 \pm 0.10	19.2 \pm 3.7	26.4 \pm 5.6	20.7 \pm 3.7	1.66 \pm 0.28	0.83 \pm 0.07	14.5 \pm 2.3	25.6 \pm 3.8	15.5 \pm 2.3
COSE	1	3.02 \pm 0.74	0.87 \pm 0.10	19.3 \pm 3.4	31.7 \pm 6.2	19.8 \pm 3.5	1.76 \pm 0.29	0.84 \pm 0.07	14.3 \pm 2.5	26.1 \pm 3.7	14.6 \pm 2.4

Table 2: Performance and training overhead of one-step generation.

Method	PESQ	ESTOI	SI-SDR	Time (ms/step)	Memory
NCSN++	2.87 \pm 0.74	0.87 \pm 0.10	19.1 \pm 3.5	70	3733
FlowSE	1.25 \pm 0.11	0.71 \pm 0.10	10.8 \pm 2.0	70	3733
MeanFlow	3.00 \pm 0.73	0.87 \pm 0.10	19.1 \pm 3.4	195	14005
COSE	3.02 \pm 0.74	0.87 \pm 0.10	19.3 \pm 3.4	112	8441

4.2. Experimental Results

Overall Performance. Table 1 summarizes COSE’s speech enhancement results on VBD and CHiME-4 datasets, compared with a range of diffusion-based and flow-matching baselines. In terms of overall performance, COSE achieves superior one-step generation: on VBD, it outperforms diffusion models such as SGMSE+ [8], StoRM [24], and VPIDM [9] even when with 15 steps, and it also surpasses advanced flow-matching methods such as FlowSE [13, 19, 30] and LARF [14]. On CHiME-4, COSE remains competitive, confirming its robustness under diverse acoustic conditions.

Table 1 also provides detailed results across different sampling steps for diffusion-based and FM baselines. It reveals that diffusion-based baselines maintain acceptable quality at 15 steps but degrade rapidly as the step count decreases, with one-step PESQ collapsing to around 1.0. Similarly, FlowSE [19], which models instantaneous velocities, fails entirely in one-step generation because the large-step discretization leads to severe trajectory deviation. In contrast, COSE consistently produces high-quality results in just one step and surpasses LARF [14]. These embody the ability of average velocity modeling to reduce cumulative errors and generate in one step.

Ablation Study. Table 2 summarizes the one-step performance and training overhead of different methods. The result shows that although MeanFlow is effective, it incurs substantial computational overhead: each training step takes 195ms and consumes 14,005MB of GPU memory, approximately 2.75x higher than that of NCSN++ and FlowSE, significantly limiting its practicality. In contrast, COSE adopts the velocity composition identity in place of the MeanFlow identity, effectively avoiding the costly *JVP* operations. This leads to

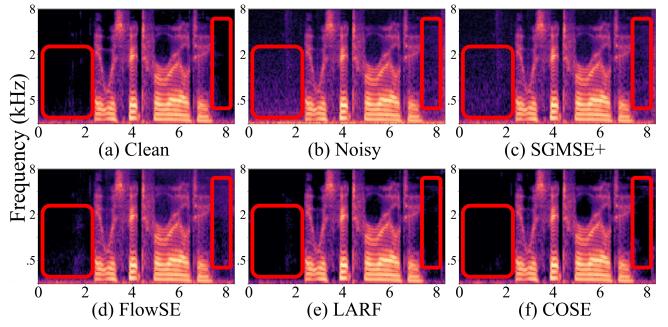


Fig. 2: Visualization of magnitude spectrum.

a notable reduction in training overhead, with training time reduced by approximately 43% and GPU memory usage by about 40%, while further improving performance. Although COSE has less training overhead than MeanFlow, its cost remains higher than NCSN++ due to two additional forward per training step.

Case Study. Figure 2 shows a visual case, where red rectangles demonstrate that COSE outperforms other baseline models in noise removal while preserving speech details. This advantage originates from COSE’s one-step generation capability, which avoids cumulative errors and improves generation quality.

5. CONCLUSION

In this work, we proposed COSE, an one-step flow matching framework that aims to guide speech enhancement process with average velocity. Our method circumvents costly *JVP* computations by leveraging a velocity composition identity. This approach significantly reduces training overhead and provides a solution equivalent to other one-step generation methods. Experiments on standard benchmarks demonstrate that COSE achieves competitive enhancement quality with significantly improved efficiency, showing the promise of one-step flow matching for practical SE applications. We hope this work encourages further development of one-step generation in SE. In future work, we plan to explore the generalization of our approach across a broader range of model architectures and datasets.

6. REFERENCES

- [1] J. S. Lim and A. V. Oppenheim, “Enhancement and bandwidth compression of noisy speech,” *Proceedings of the IEEE*, vol. 67, no. 12, pp. 1586–1604, 2005.
- [2] B. Li, A. Gulati, J. Yu, T. N. Sainath, C.-C. Chiu, A. Narayanan, S.-Y. Chang, R. Pang, Y. He, J. Qin, et al., “A better and faster end-to-end model for streaming ASR,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. IEEE, 2021, pp. 5634–5638.
- [3] C. Zheng, H. Zhang, W. Liu, X. Luo, A. Li, X. Li, and B. C. J. Moore, “Sixty years of frequency-domain monaural speech enhancement: From traditional to deep learning methods,” *Trends in Hearing*, vol. 27, 2023.
- [4] Y. Tsao, P. Plantinga, M. Ravanelli, S.-W. Fu, C. Yu, T.-A. Hsieh, and X. Lu, “MetricGAN+: An improved version of MetricGAN for speech enhancement,” *Interspeech*, 2021.
- [5] B. Kang, X. Zhu, Z. Zhang, Z. Ye, M. Liu, Z. Wang, Y. Zhu, G. Ma, J. Chen, L. Xiao, C. Weng, W. Xue, and L. Xie, “LLaSE-G1: Incentivizing generalization capability for LLaMA-based speech enhancement,” in *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [6] Y.-J. Lu, Z.-Q. Wang, S. Watanabe, A. Richard, C. Yu, and Y. Tsao, “Conditional diffusion probabilistic model for speech enhancement,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. IEEE, 2022, pp. 1–5.
- [7] H. Yen, F. G. Germain, G. Wichern, and J. Le Roux, “Cold diffusion for speech enhancement,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. IEEE, 2023, pp. 1–5.
- [8] J. Richter, S. Welker, J.-M. Lemercier, B. Lay, and T. Gerkmann, “Speech enhancement and dereverberation with diffusion-based generative models,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 31, pp. 2351–2364, 2023.
- [9] Z. Guo, Q. Wang, J. Du, J. Pan, Q.-F. Liu, and C.-H. Lee, “A variance-preserving interpolation approach for diffusion models with applications to single channel speech enhancement and recognition,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 2024.
- [10] P. Gonzalez, Z.-H. Tan, J. Østergaard, J. Jensen, T. S. Alstrøm, and T. May, “Investigating the design space of diffusion models for speech enhancement,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 2024.
- [11] W. Tai, Y. Lei, F. Zhou, G. Trajcevski, and T. Zhong, “DOSE: Diffusion dropout with adaptive prior for speech enhancement,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024, vol. 36.
- [12] R. Korostik, R. Nasretdinov, and A. Jukić, “Modifying flow matching for generative speech enhancement,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. IEEE, 2025, pp. 1–5.
- [13] S. Lee, S. Cheong, S. Han, and J. W. Shin, “FlowSE: Flow matching-based speech enhancement,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. IEEE, 2025, pp. 1–5.
- [14] Z. Li and N. Inoue, “Locally aligned rectified flow model for speech enhancement towards single-step diffusion,” in *Proc. Interspeech*, 2024, pp. 2195–2199.
- [15] K. Frans, D. Hafner, S. Levine, and P. Abbeel, “One step diffusion via shortcut models,” in *International Conference on Learning Representations (ICLR)*, 2025.
- [16] Z. Geng, M. Deng, X. Bai, J. Z. Kolter, and K. He, “Mean flows for one-step generative modeling,” *arXiv preprint*, 2025.
- [17] L. Yang, Z. Zhang, Z. Zhang, X. Liu, M. Xu, W. Zhang, C. Meng, S. Ermon, and B. Cui, “Consistency flow matching: Defining straight flows with velocity consistency,” *arXiv preprint*, 2024.
- [18] X. Liu, C. Gong, and Q. Liu, “Flow straight and fast: Learning to generate and transfer data with rectified flow,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [19] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le, “Flow matching for generative modeling,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [20] A. Paszke, S. Gross, S. Chintala, G. Chanan, E. Yang, Z. DeVito, Z. Lin, A. Desmaison, L. Antiga, and A. Lerer, “Automatic differentiation in pytorch,” in *NIPS 2017 Autodiff Workshop*, 2017.
- [21] A. G. Baydin, B. A. Pearlmutter, A. A. Radul, and J. M. Siskind, “Automatic differentiation in machine learning: a survey,” *Journal of machine learning research*, pp. 1–43, 2018.
- [22] L. Li, S. Hurault, and J. M. Solomon, “Self-consistent velocity matching of probability flows,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023, vol. 36, pp. 57038–57057.
- [23] G. Teschl, *Ordinary differential equations and dynamical systems*, vol. 140, American Mathematical Soc., 2012.
- [24] J.-M. Lemercier, J. Richter, S. Welker, and T. Gerkmann, “StoRM: A diffusion-based stochastic regeneration model for speech enhancement and dereverberation,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 31, pp. 2724–2737, 2023.
- [25] C. Veaux, J. Yamagishi, and S. King, “The voice bank corpus: Design, collection and data analysis of a large regional accent speech database,” in *International Conference Oriental COCOSDA*. IEEE, 2013, pp. 1–4.
- [26] E. Vincent, S. Watanabe, A. A. Nugraha, J. Barker, and R. Marxer, “An analysis of environment, microphone and data simulation mismatches in robust speech recognition,” *Computer Speech & Language*, vol. 46, pp. 535–557, 2017.
- [27] A. W. Rix, J. G. Beerends, M. P. Hollier, and A. P. Hekstra, “Perceptual evaluation of speech quality (PESQ) - a new method for speech quality assessment of telephone networks and codecs,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, 2001, pp. 749–752.
- [28] J. Jensen and C. H. Taal, “An algorithm for predicting the intelligibility of speech masked by modulated noise maskers,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 24, no. 11, pp. 2009–2022, 2016.
- [29] J. Le Roux, S. Wisdom, H. Erdogan, and J. R. Hershey, “SDR-half-baked or well done?,” in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*. IEEE, 2019, pp. 626–630.
- [30] Z. Wang, Z. Liu, X. Zhu, Y. Zhu, M. Liu, J. Chen, L. Xiao, C. Weng, and L. Xie, “FlowSE: Efficient and high-quality speech enhancement via flow matching,” in *Proc. Interspeech*, 2025, pp. 4858–4862.